

EECE 5550 Mobile Robotics

Lecture 14: Motion Planning

Derya Aksaray

Assistant Professor

Department of Electrical and Computer Engineering



Northeastern
University

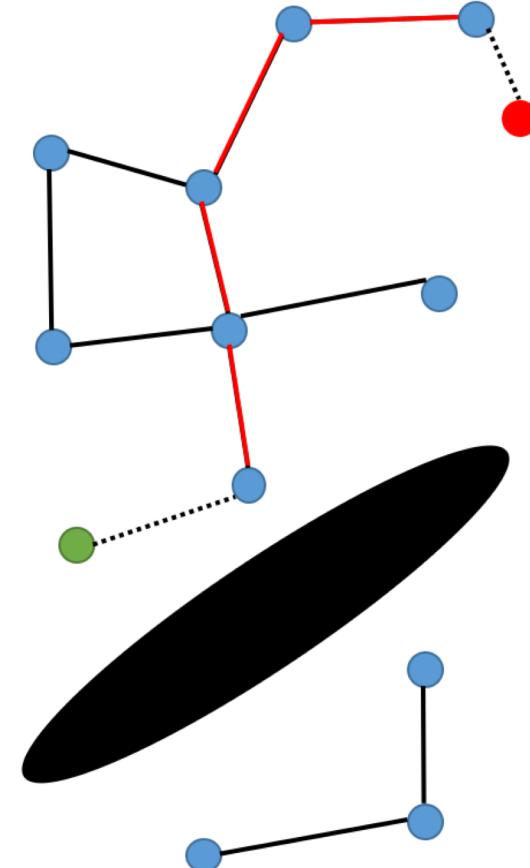
Plan of the day

Last time: Planning as search

- Definition of planning problems & solutions
- Planning as graph search
- Graph search algorithms

Today: Application to robot motion planning

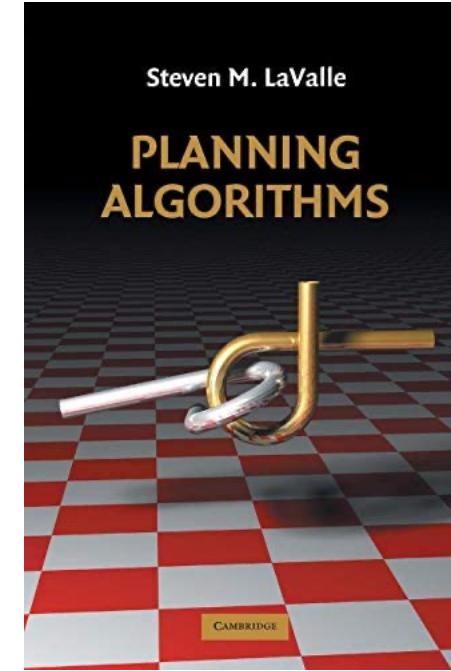
- Robot workspaces & configuration spaces
- Grid-based motion planners
- Sampling-based planners for **high-dimensional** spaces
 - Probabilistic road maps (**PRMs**)
 - Rapidly-exploring random trees (**RRTs**)



References



Lecture “Planning II” from ETH Zurich’s
Autonomous Mobile Robots course



Chapters 4,5

Classic papers:

- “Probabilistic Roadmaps for Path Planning in High-Dimensional Configuration Spaces”
- “Randomized Kinodynamic Planning”

Classical goals of motion planning

Assumptions

- Static world
- Robot knows about the map

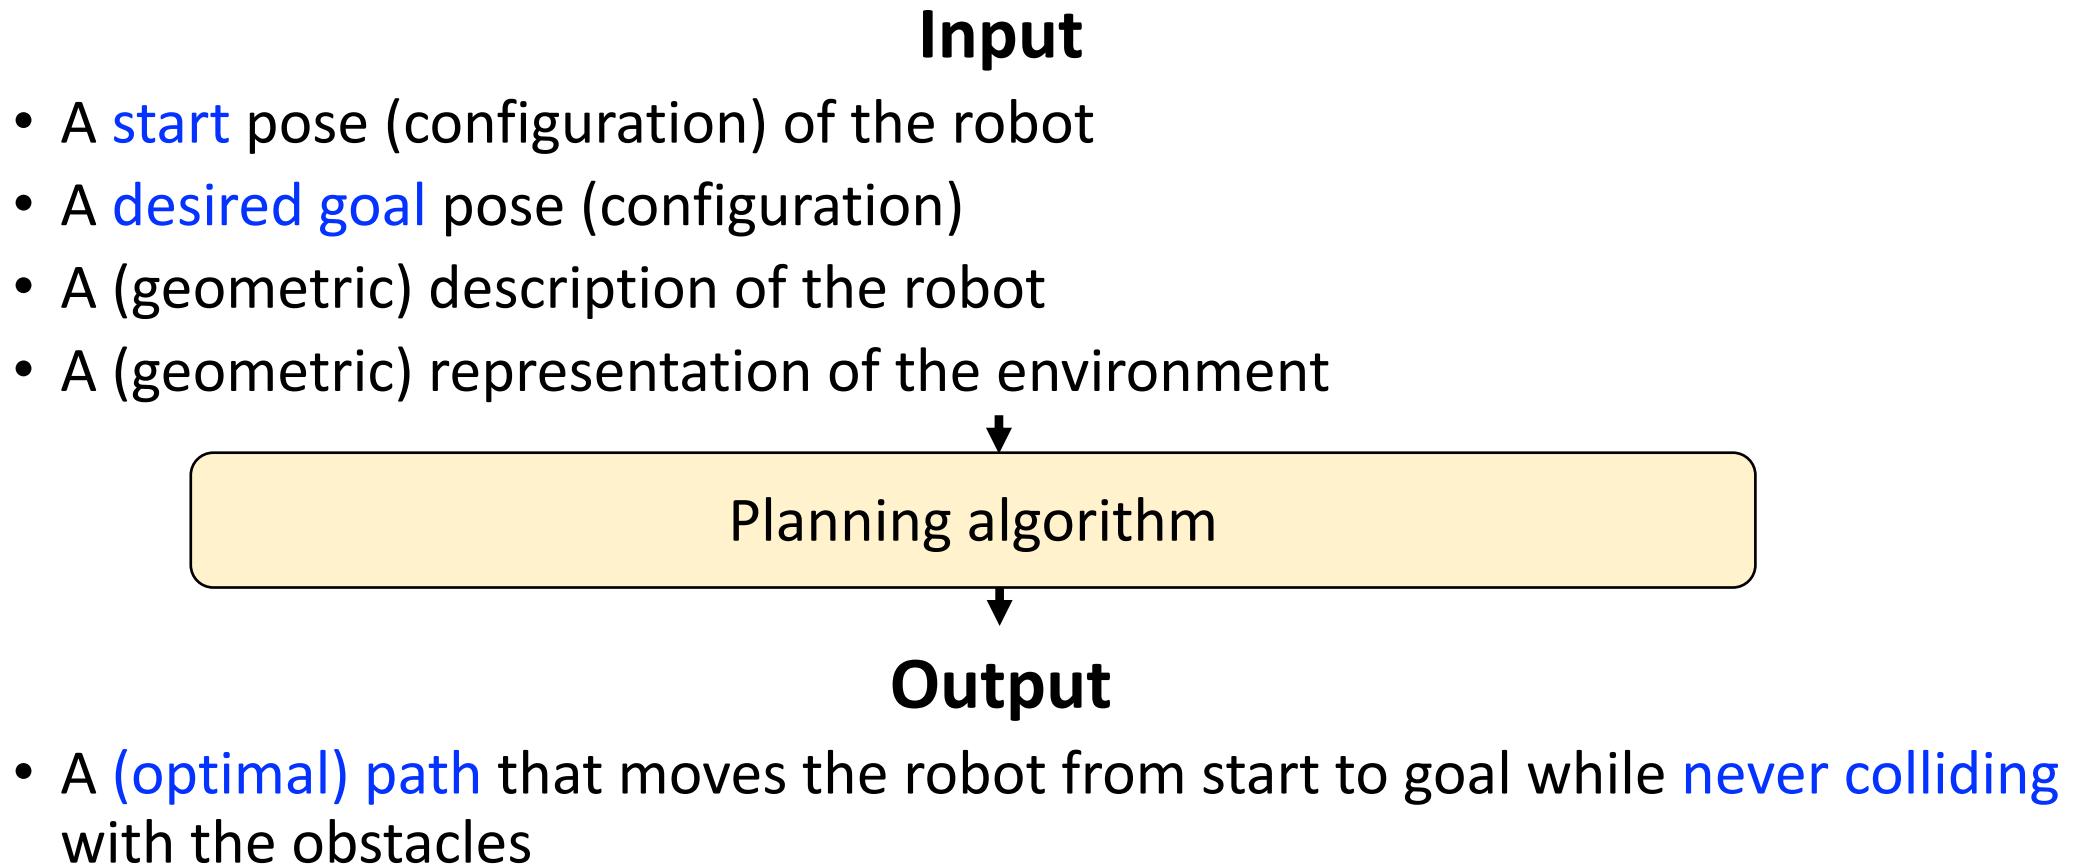
Goals

- Collision-free trajectories
- Robot should reach the goal location with min cost
(e.g., min time, min distance)

Planning in dynamic environments (not the scope of this course)

- More suitable to real-world applications
- Would require some knowledge about the movement of the surroundings

Motion planning problem



Workspace and configuration space

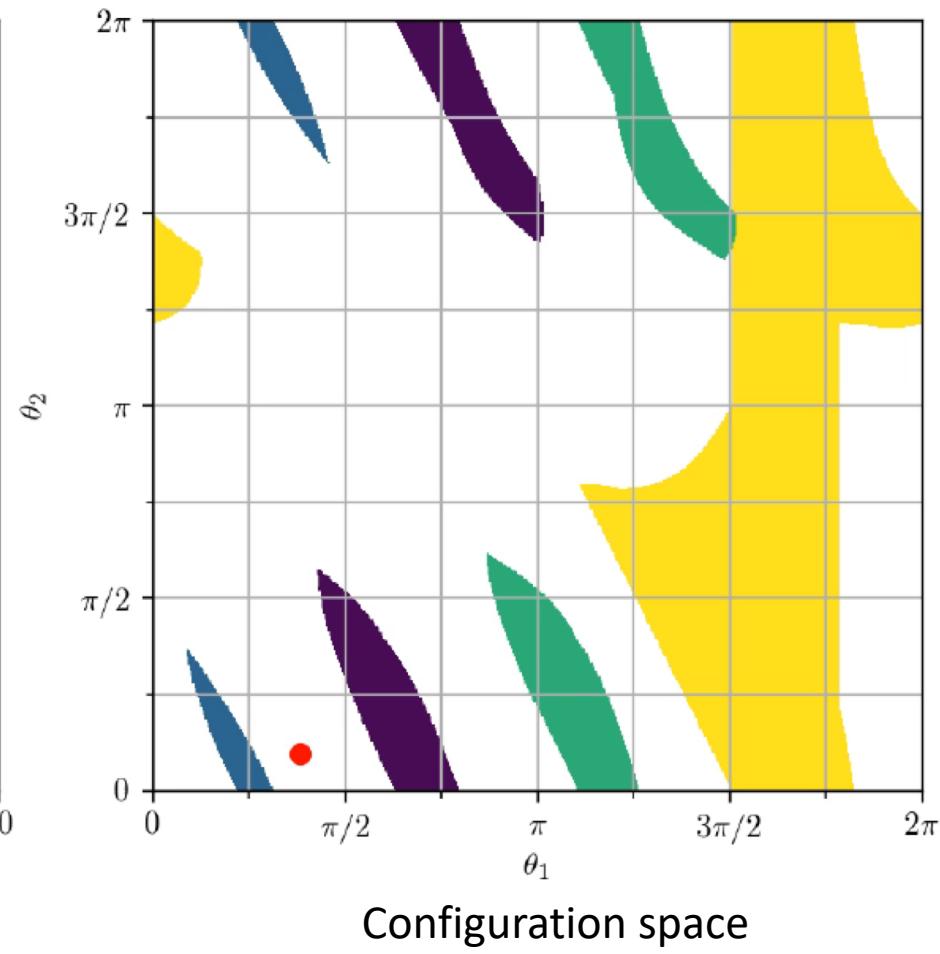
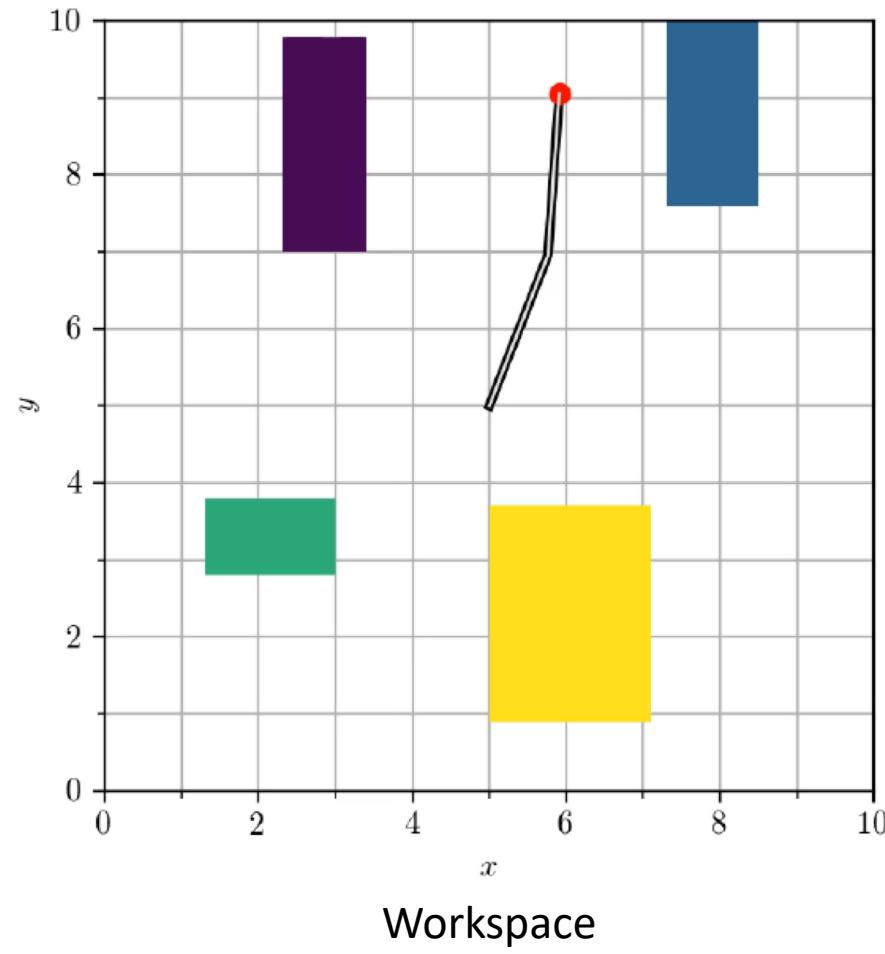
Workspace: The **environment** in which the robot is operating.

Often modeled at the level of **free** and **occupied** space

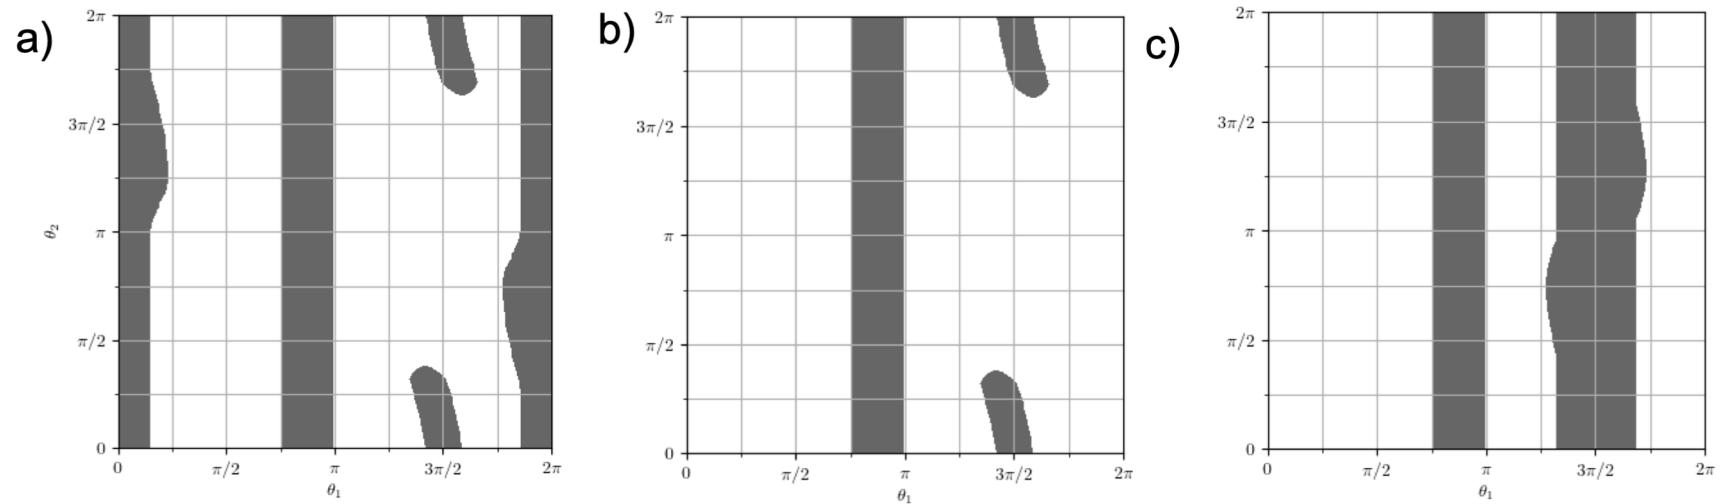
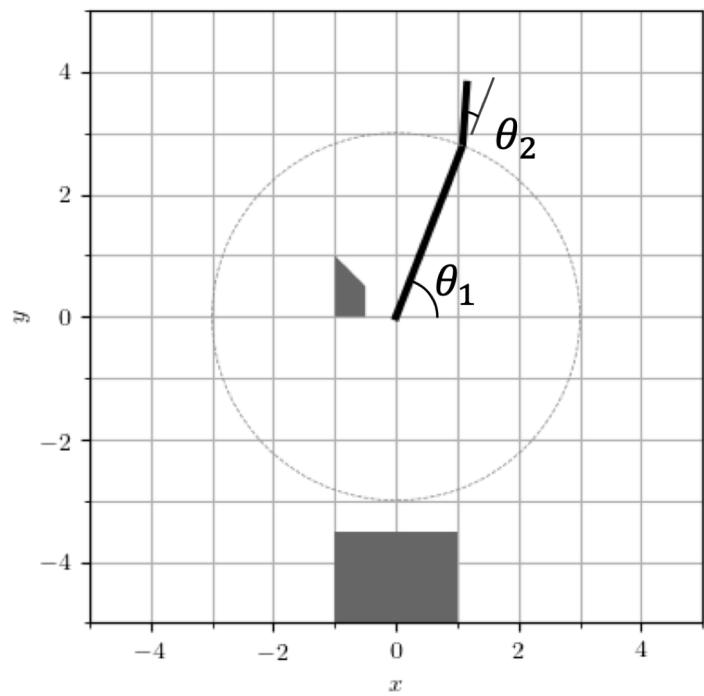
Configuration space (C-space): The set of all possible **robot configurations (states)**

- A configuration of a robot is a minimal set of variables that specifies the position and orientation of each rigid body composing the robot.
- The number of **degrees of freedom** of a robot is the dimension of the C-space.
 - The min # of variables required to fully specify the position and orientation of each rigid body belonging to the robot.
- Partially determined by the workspace (no collisions permitted).

Workspace and configuration space: two-link arm



Exercise: configuration space



Some robot examples: C-space

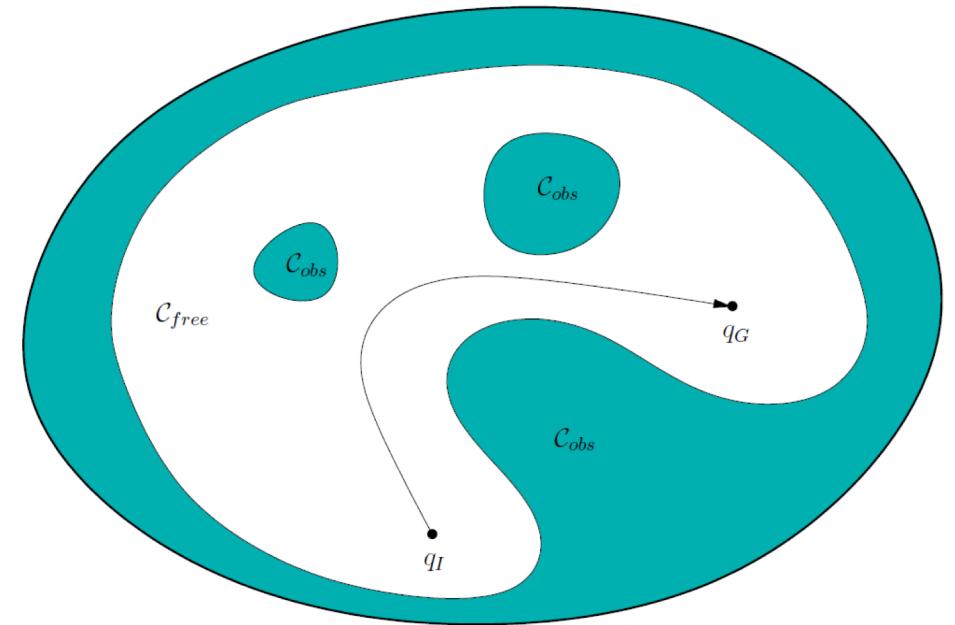
- Point robot in \mathbb{R}^2 - Suppose that the robot can only translate through space
 - Robot's configuration: (x, y)
 - C-space: 2D plane
- 2D robot that can translate and rotate in a 2D workspace
 - Robot's configuration: (x, y, θ)
 - C-space: $SE(2) = \mathbb{R}^2 \times SO(2)$
- A space robot moving in 3D space
 - Robot's configuration: $(x, y, z, \theta, \phi, \psi)$
 - C-space: $SE(3) = \mathbb{R}^3 \times SO(3)$
- A fixed robot with n arm linkages
 - Robot's configuration: n angles
 - C-space: n-dimensional space

Robot motion planning

Given:

- Workspace W , partitioned into *free* and *occupied* subsets
- Robot *configuration space* C , partitioned into corresponding *free* and *occupied* subsets
- Initial configuration $q_I \in C_{free}$
- Goal configuration $q_G \in C_{free}$

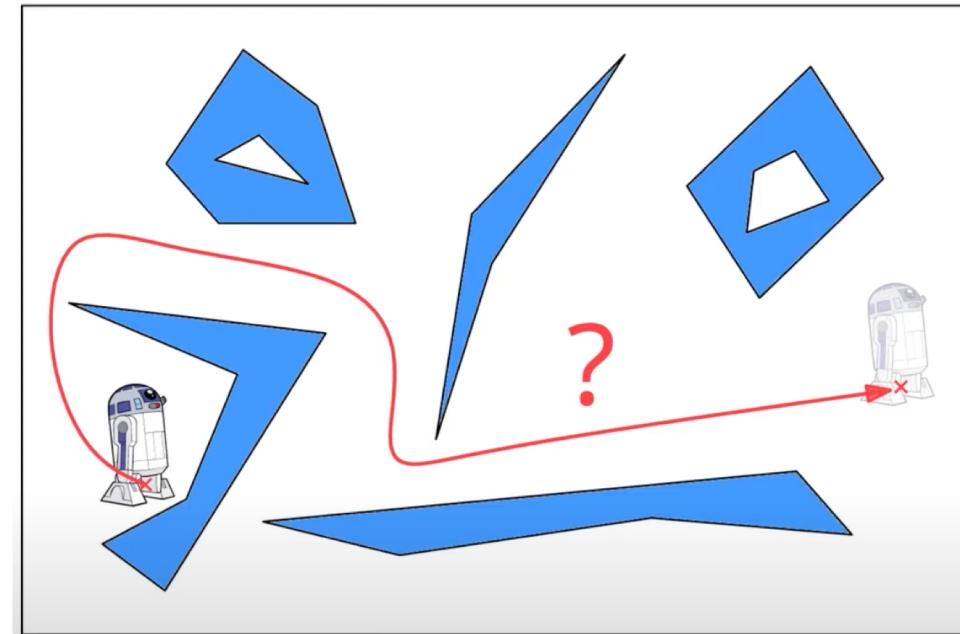
Find: A path $\tau: [0,1] \rightarrow C_{free}$ such that $\tau(0) = q_I$ and $\tau(1) = q_G$



Configuration space

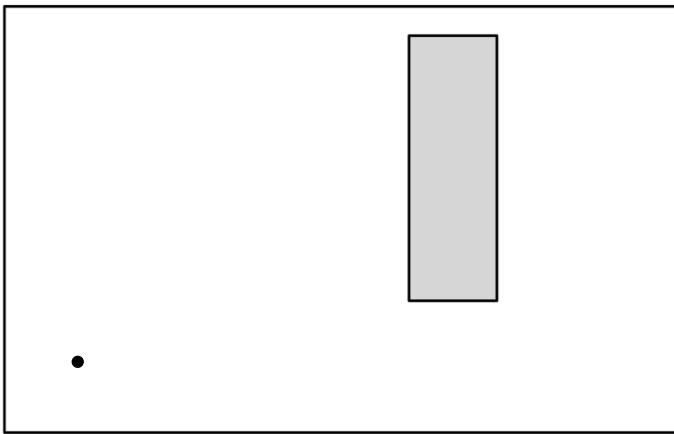
What are critical when constructing free and occupied c-space?

- Shape of robot
 - (point shaped)
- Shape of obstacle
 - (polygonal)
- Motion (e.g., rotation) of robot
 - (only translation)
- Motion of obstacles
 - (no motion)
- 2D, 3D
 - (2D)



Workspace and configuration space of a robot

Simplest case

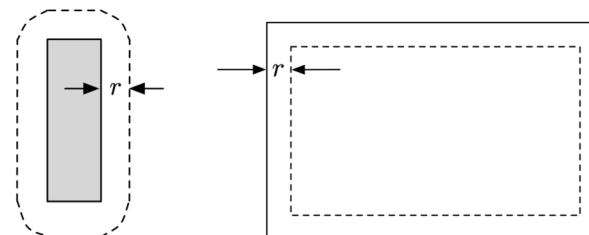
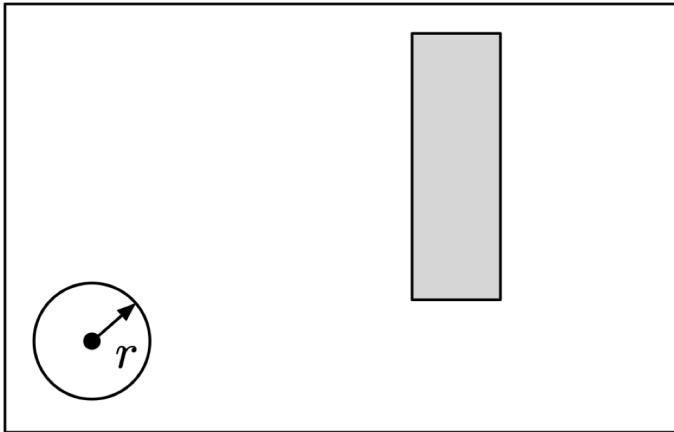


Workspace: Rectangular 2D region with an obstacle

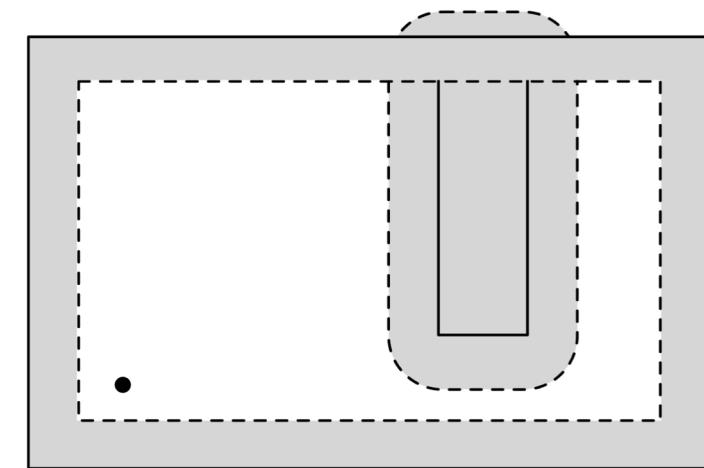
Robot: Point robot

Free configuration space: The white region

Robot: Disc robot

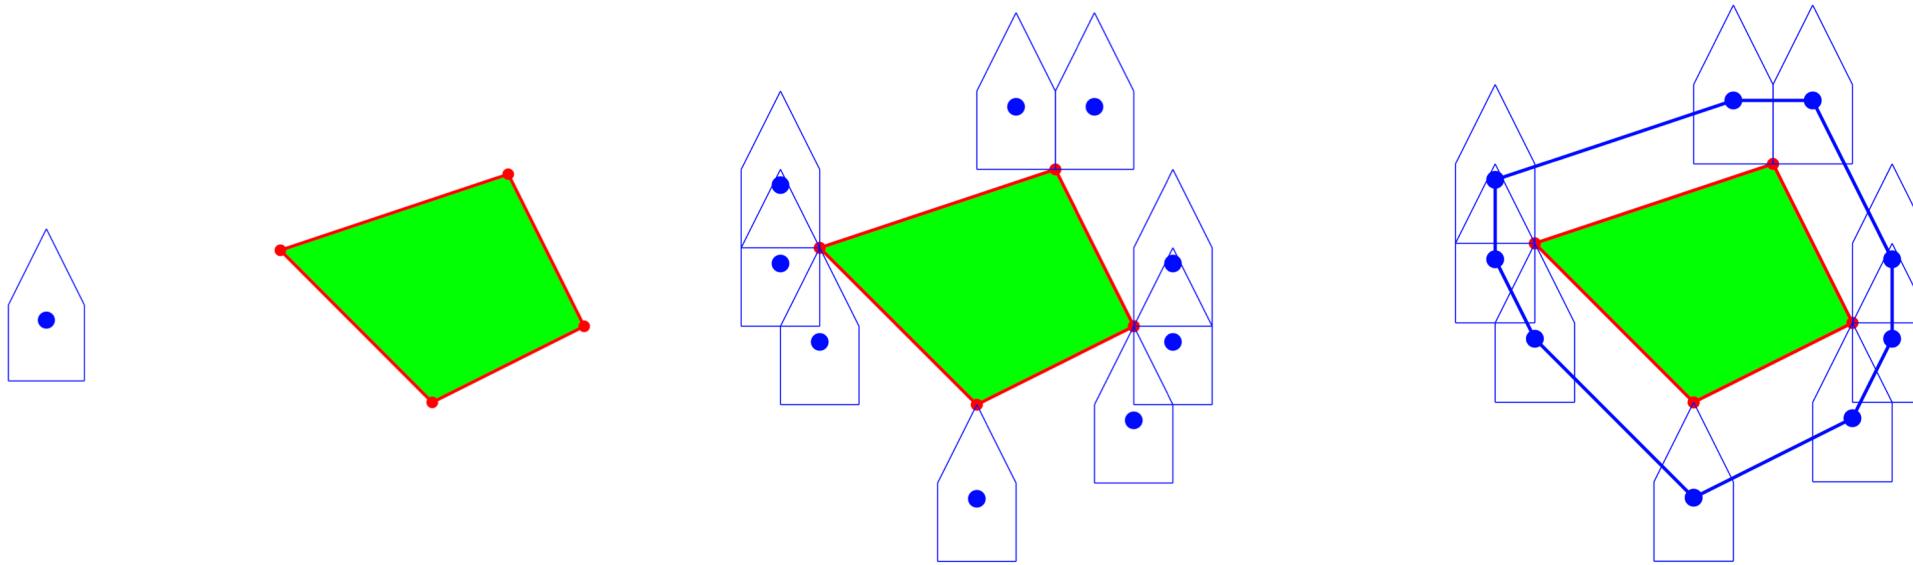


Expand all obstacles
by radius r



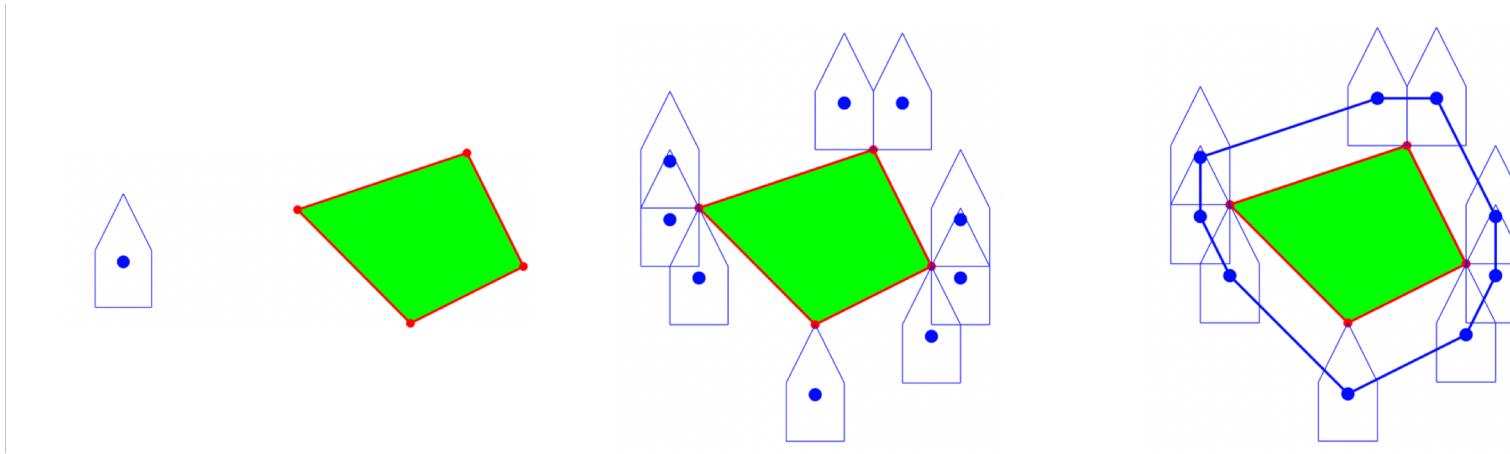
Free configuration space of non-circular robot

Example: Consider a polygonal robot R which will navigate around a polygon obstacle P with planar translation.



Free configuration space of non-circular robot

- How do we formally grow the obstacles by the robot's shape so that we can transform the motion planning problem into one where the robot is treated as a point?



Minkowski sum: Let P and Q be sets of points. The Minkowski sum $P \oplus Q$ of P and Q is

$$P \oplus Q = \{ p + q \mid p \in P, q \in Q \},$$

where $p+q=(p_x + q_x, p_y + q_y)$, for $p=(p_x, p_y)$ and $q=(q_x, q_y)$.

- summing all vertices of P and Q into the set S ,
- taking the convex hull of the points in S .

Robot motion planning: Challenges

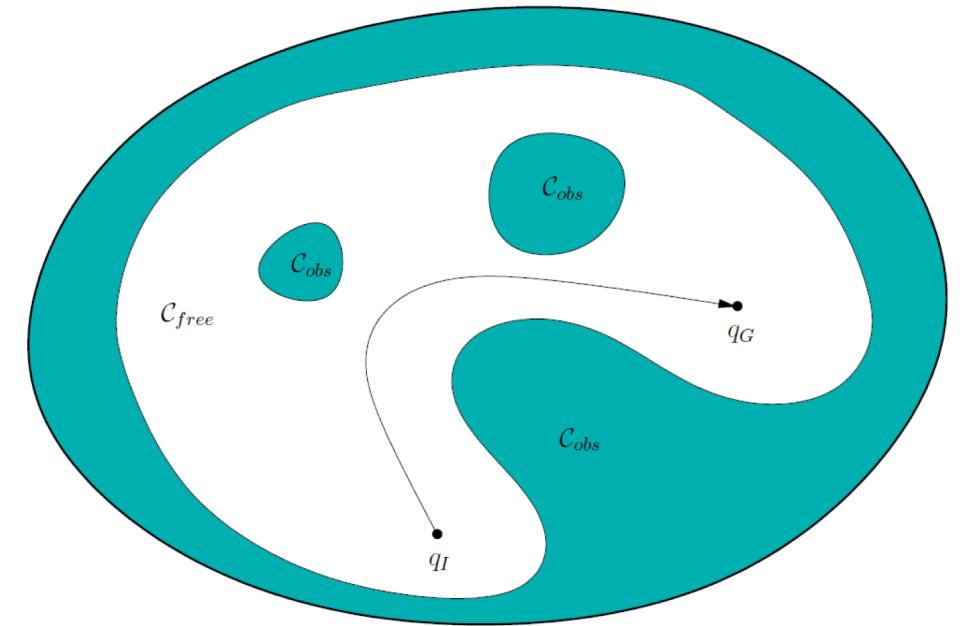
Basic challenge: Workspaces and robot configuration spaces are generally *continuous*

⇒ Very hard to model arbitrary continuous shapes

One natural approach: *Discretize* the C-space to get a *finite approximation*, and then solve the problem via *search*.

Two general approaches to discretize C-spaces:

- **Combinatorial planning:** Characterize the free C-space by explicitly capturing the connectivity of the free space into a graph, then graph search
- **Sampling-based planning:** Incrementally search C-space with the consideration of collision detection



Some important definitions for search alg.

A path planning algorithm is:

complete:

- If a path exists, it finds it in *finite* time
- If a path does not exist, it returns in *finite* time

probabilistically complete:

- if a path exists, the probability of finding it tends to go to 1 as the iterations go to infinity

resolution complete:

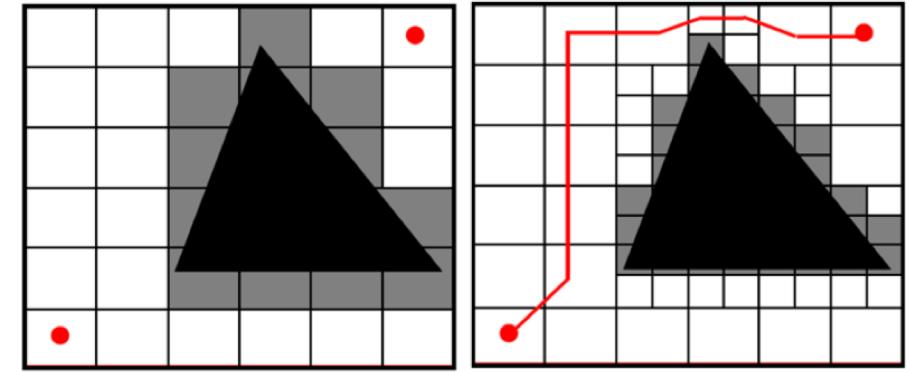
- If a path exists, it finds one; otherwise, it terminates and reports no solution within a specified resolution

sound:

- if guaranteed to never cross an obstacle

optimal:

- if guaranteed to find the shortest path (if it exists)



Combinatorial planning

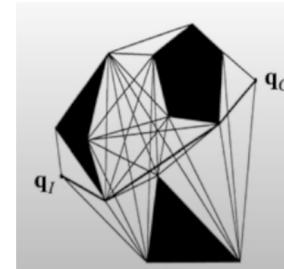
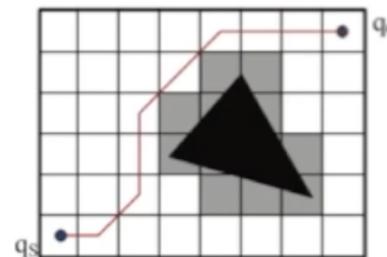
Step 1: Compute C-space

Step 2: Generate a roadmap (i.e., a graph) in free C-space

- Cell decomposition methods, visibility graphs, occupancy grid maps, etc.

Step 3: Compute the min cost path from initial configuration to the goal configuration

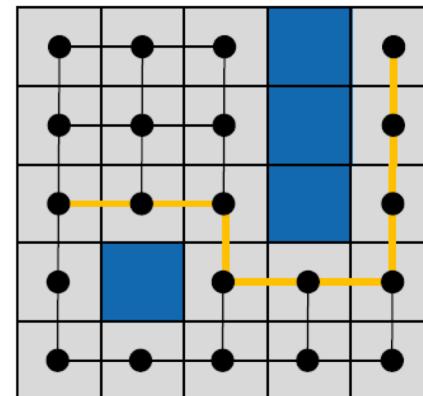
- Cast as a graph search algorithm
- Dijkstra, A*, etc.



Simple case: route planning for a planar robot

Recall:

- **Occupancy grids** provide a discrete model of free and occupied space
- For a *point* robot, C_{free} is just the set of unoccupied cells
- For *disc* robot, can also “grow” occupied cells to account for the body



Now: Construct a graph $G = (V, E)$ where:

- Vertices are free cells
- Edges model connectivity between free cells

Then: Motion planning is just graph search!

Grid-based representations

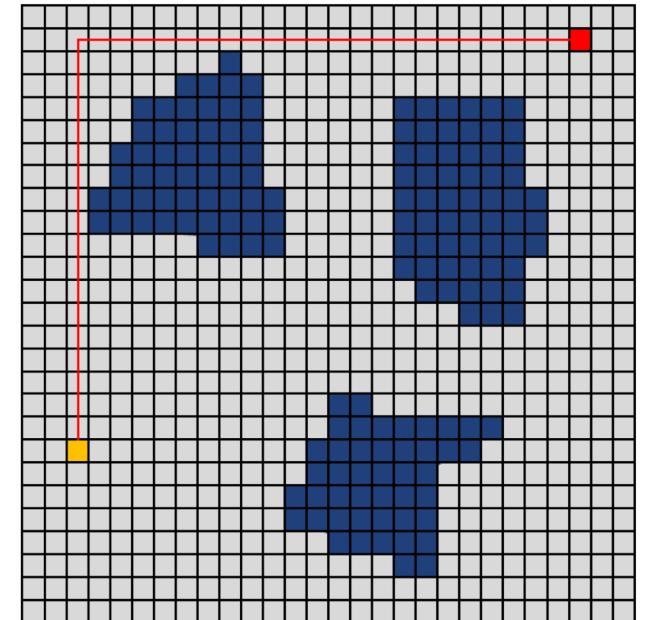
More generally, one can “voxelize” any *generic* configuration space C_{free} , and then apply the same graph search strategy

Pros:

- Very simple idea
- Easy to implement
- **Resolution completeness**: If there is a feasible plan, this approach is *guaranteed* to find it as grid resolution $r \rightarrow 0$.

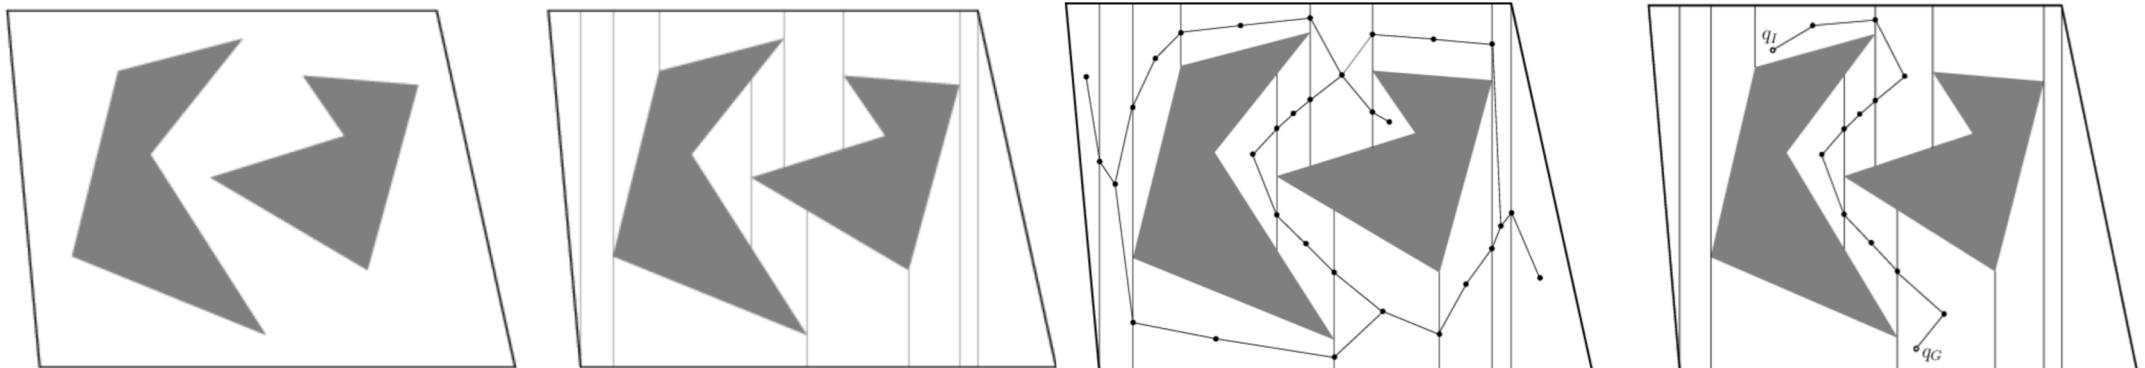
Con:

- Voxelization can introduce weird artifacts
- Not always clear how to choose resolution r : How small is “small enough”??
- **Curse of dimensionality**: Number of cells N in the voxel grid grows as $O(r^d)$!



Vertical cell decomposition

- Divides the space into non-overlapping cells.
- Add vertical line segment to every obstacles' vertices.
- Create graph (roadmap): vertices in trapezoids and vertical segments



- If there is a path from start to goal, such a decomposition and graph search over that would guarantee to find the solution.
- Complex shape obstacles can lead to graphs that can not result in the actual shortest path.

Combinatorial planning - challenges

Exact method (no approximation needed) and complete but,

C-space has high dimensionality

- A rigid body in 3D space has a C-space with 6 dimensions ($C = \mathbb{R}^3 \times SO(3)$)

Simple obstacles have complex C-obstacles

- Impractical to compute explicit representation of free space for high dof robots

More attention on sampling-based planning...

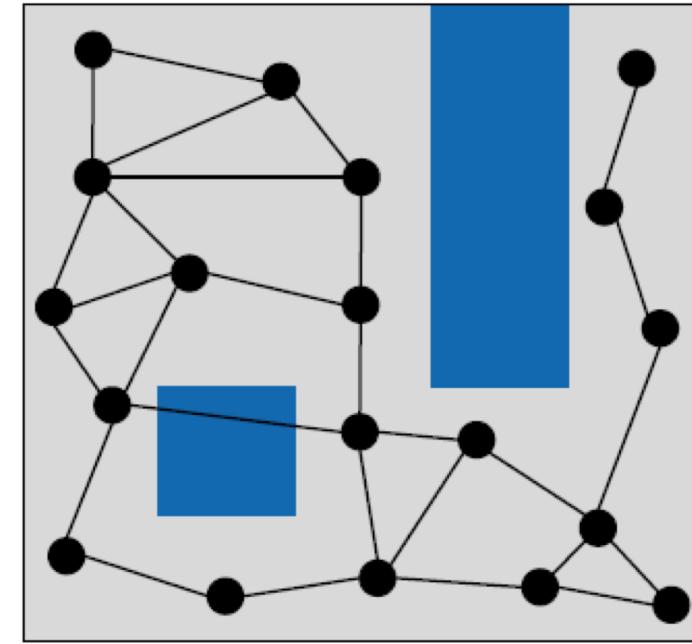
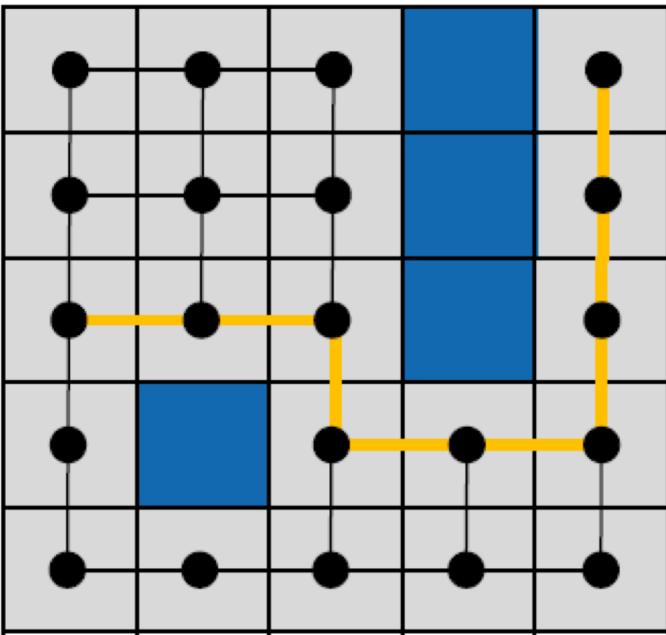
Sampling-based planning

Main idea: Rather than trying to capture *every* point in the configuration space C via voxelization, let's *randomly sample* a *representative set S* of points in C_{free}

Then:

- $S \subset C_{free}$ provides a *sample-based inner approximation* of C_{free}
- **NB:** It's often easy to plan feasible motions between two *nearby* points $x, y \in S$
(E.g.: a straight-line path often suffices ...)
- If we draw an edge e between two points $x, y \in S$ whenever we can *locally* plan a feasible path from one to the other, we get a graph $G = (V, E)$ that models *reachability*

Grid- vs. sample-based planners



Sampling-based planning: General framework

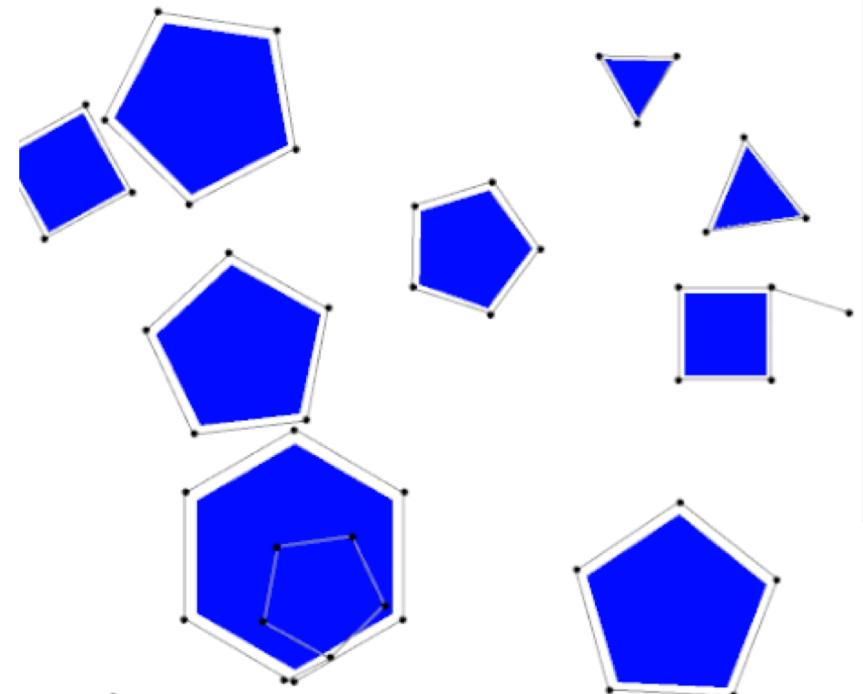
Starting with an empty graph G , repeat:

1. Sample a random point $x \in C_{free}$ and add to G
[Q: How to sample x ?]
2. For each vertex $y \in G$ s.t. $d(x, y) < \epsilon$, try to **plan a path from x to y**

NB: This requires:

- A suitable notion of **distance** for C
- A **fast local planner**

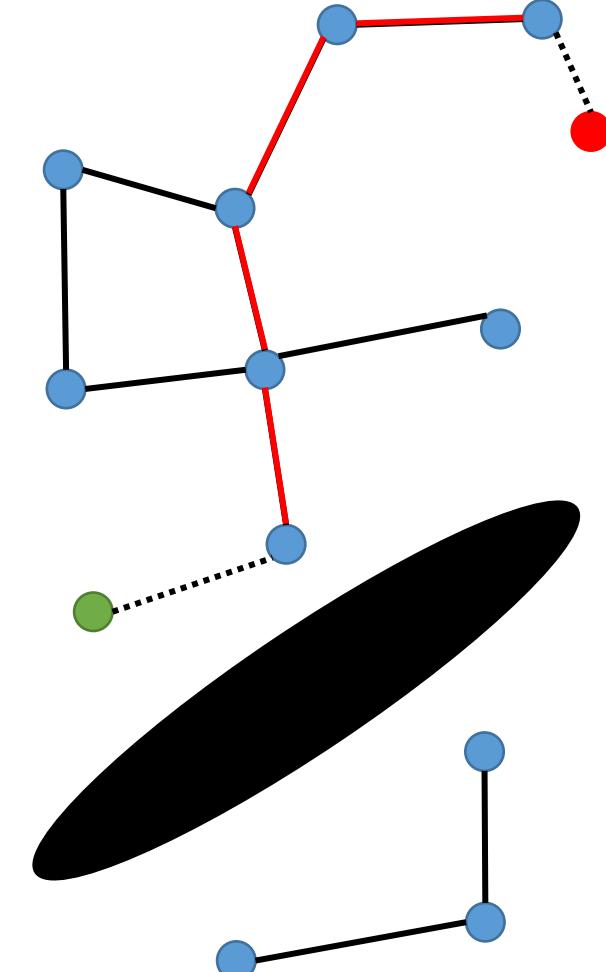
3. If a feasible path is found, add edge (x, y) to G .



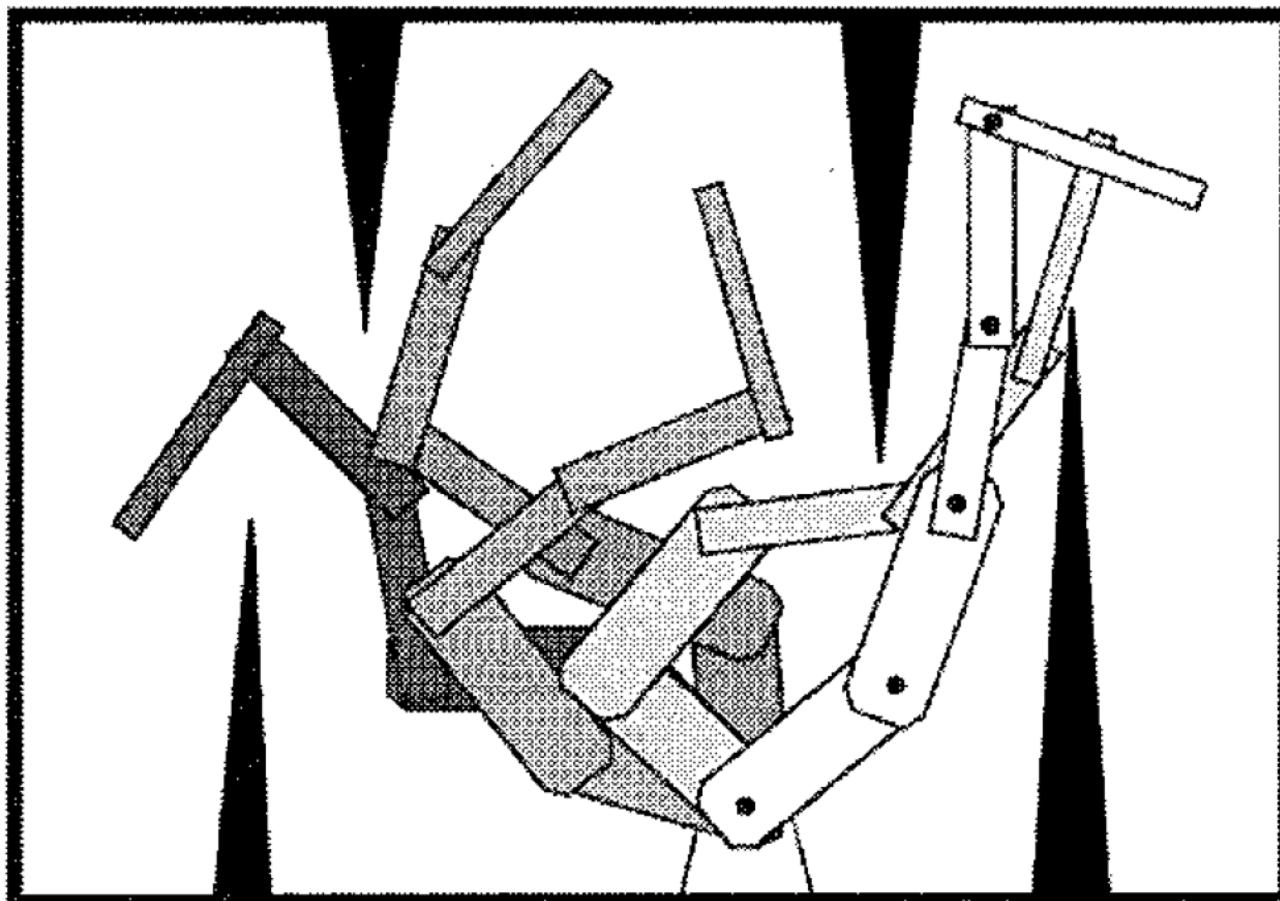
The Probabilistic Roadmap (PRM) Algorithm

Two-phase algorithm for sampling-based planning:

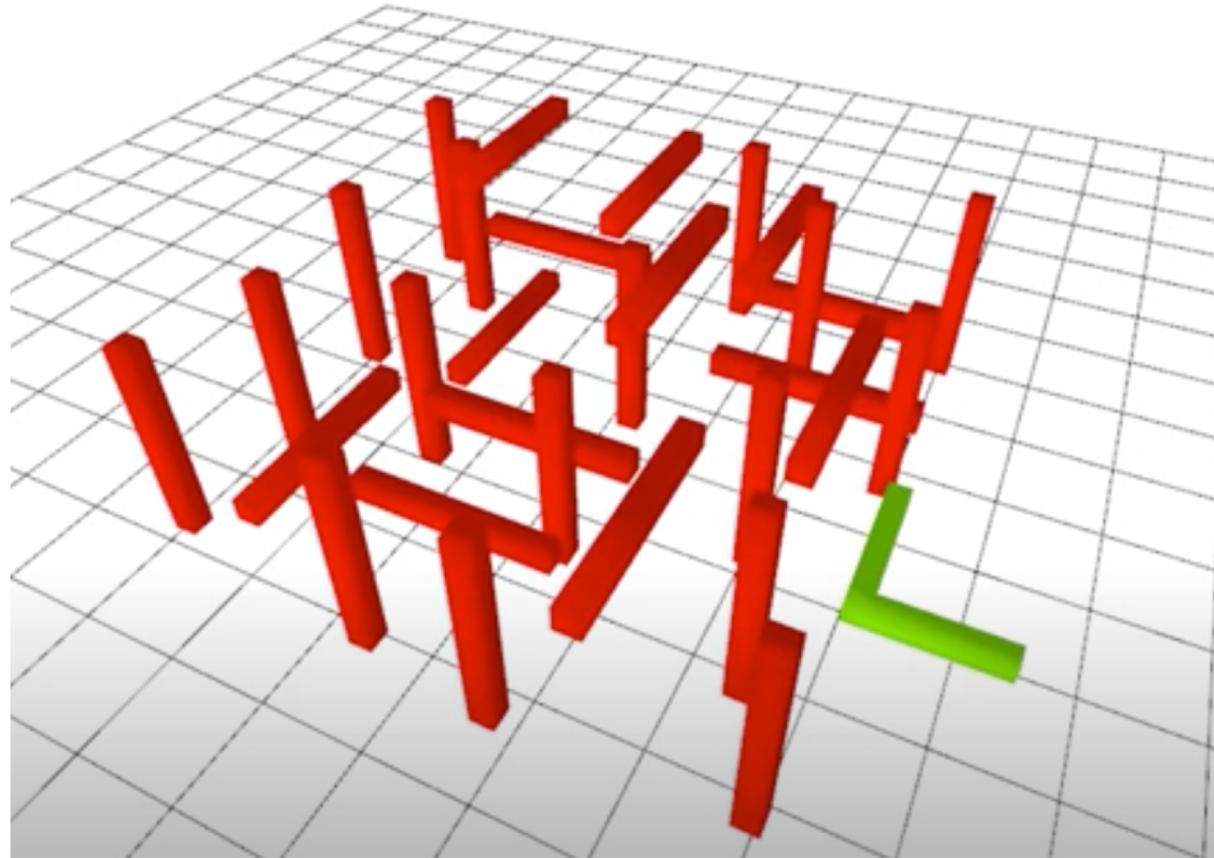
1. **Construction:** Build a **roadmap** (graph) by randomly sampling points over the **entire** configuration space and planning local paths
2. **Query:** At run-time, given initial point $x \in G$ and goal $y \in G$:
 - Plan **local paths** from x and y to (nearby) vertices in the **same connected component** of G
 - Graph search!



Example: PRM planning with a 5 DOF robotic arm



Example: Solving the piano mover's problem with PRMs



https://www.youtube.com/watch?v=3_S3GPxAMYA

Probabilistic Roadmaps: Key Properties

Recap: Builds a roadmap (graph) G over the *entire* space C_{free} by joining a *sparse sample set* using *local planning*

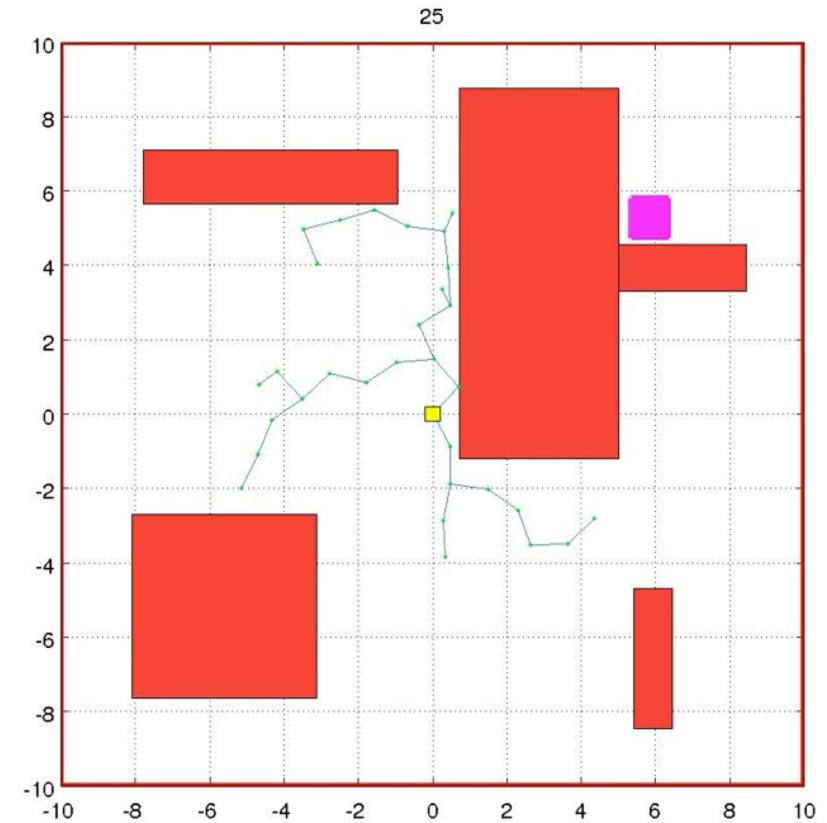
- **Much** more scalable than grid methods
 - We *sparsely (inner) approximate* true space C_{free} using sample-set S
 - Works well for high degree-of-freedom robots (e.g. robot arms, etc.)
- **Probabilistically complete**: If a feasible plan exists, the probability of finding it approaches 1 as the number of samples grows.
- Convenient “anytime” flavor: We can stop building the graph whenever we want, and we still get something useful
- Reusable! Enables fast *multi-query* planning

BUT: Costs a lot up-front (we build the roadmap G over the *entire* space C_{free})

Q: Can we get something faster in the *single-query* case?

Rapidly-exploring Random Trees (RRTs)

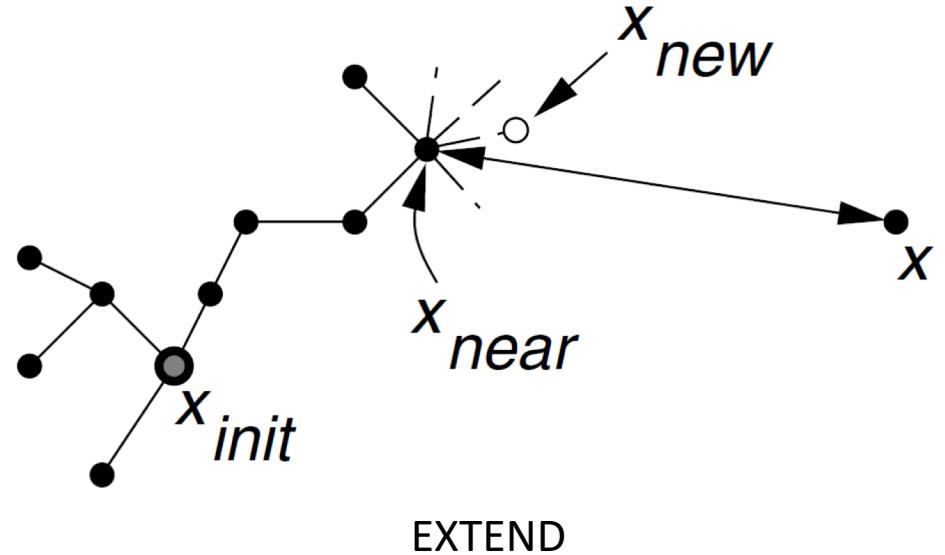
Main idea: Instead of building a *graph* by sampling vertices *uniformly* over C_{free} , we **build a tree outwards from the initial state x towards the goal y .**



RRT algorithm

```
BUILD_RRT( $x_{init}$ )
1    $\mathcal{T}_{init}(x_{init})$ ;
2   for  $k = 1$  to  $K$  do
3        $x_{rand} \leftarrow \text{RANDOM\_STATE}()$ ;
4       EXTEND( $\mathcal{T}$ ,  $x_{rand}$ );
5   Return  $\mathcal{T}$ 
```

```
EXTEND( $\mathcal{T}$ ,  $x$ )
1    $x_{near} \leftarrow \text{NEAREST\_NEIGHBOR}(x, \mathcal{T})$ ;
2   if NEW\_STATE( $x$ ,  $x_{near}$ ,  $x_{new}$ ,  $u_{new}$ ) then
3        $\mathcal{T}.\text{add\_vertex}(x_{new})$ ;
4        $\mathcal{T}.\text{add\_edge}(x_{near}, x_{new}, u_{new})$ ;
5       if  $x_{new} = x$  then
6           Return Reached;
7       else
8           Return Advanced;
9   Return Trapped;
```



NEW_STATE: x_{new} is gotten from x_{near} and u_{new} by
forward simulation using system dynamics:

$$\dot{x} = f(x, u)$$

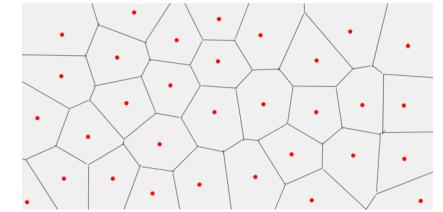
Fig. 5. Basic rapidly exploring random tree construction algorithm.

Key point: Unlike PRM, in RRT we *do not* require that the control u_{new} drives x_{near} to x ; only that it *makes progress towards* x .

What makes RRTs “rapidly exploring”?

Voronoi diagram:

- Each point decomposes the space around it into some region of influence
- All the points inside a Voronoi partition is closest to the point generated that partition

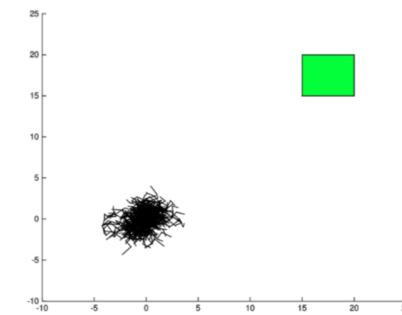


RRT does not have a particular step to create Voronoi diagram; however, one can see a Voronoi interpretation of the algorithm.

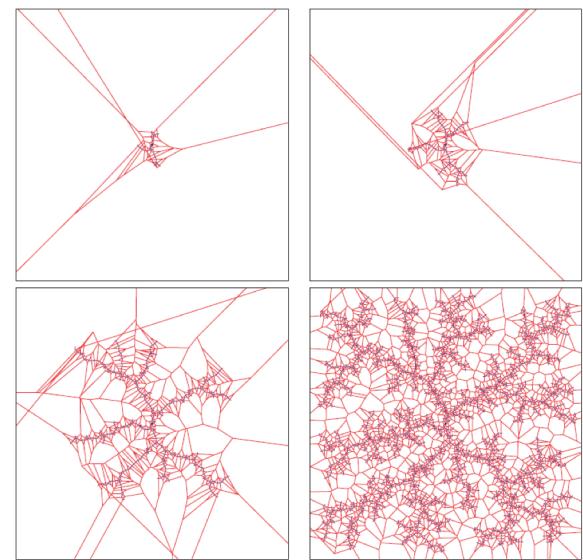
Voronoi bias:

- Samples are selected uniformly over the space.
- Such random sampling tends to place new vertices in larger Voronoi cells (\Rightarrow unexplored regions)
- This helps to expand the tree to rapidly cover the space.
- Generating new points in the space, and extending from there better than naive sampling

Naive sampling



Pick a node on the tree
Extend from there



RRT: Key Properties

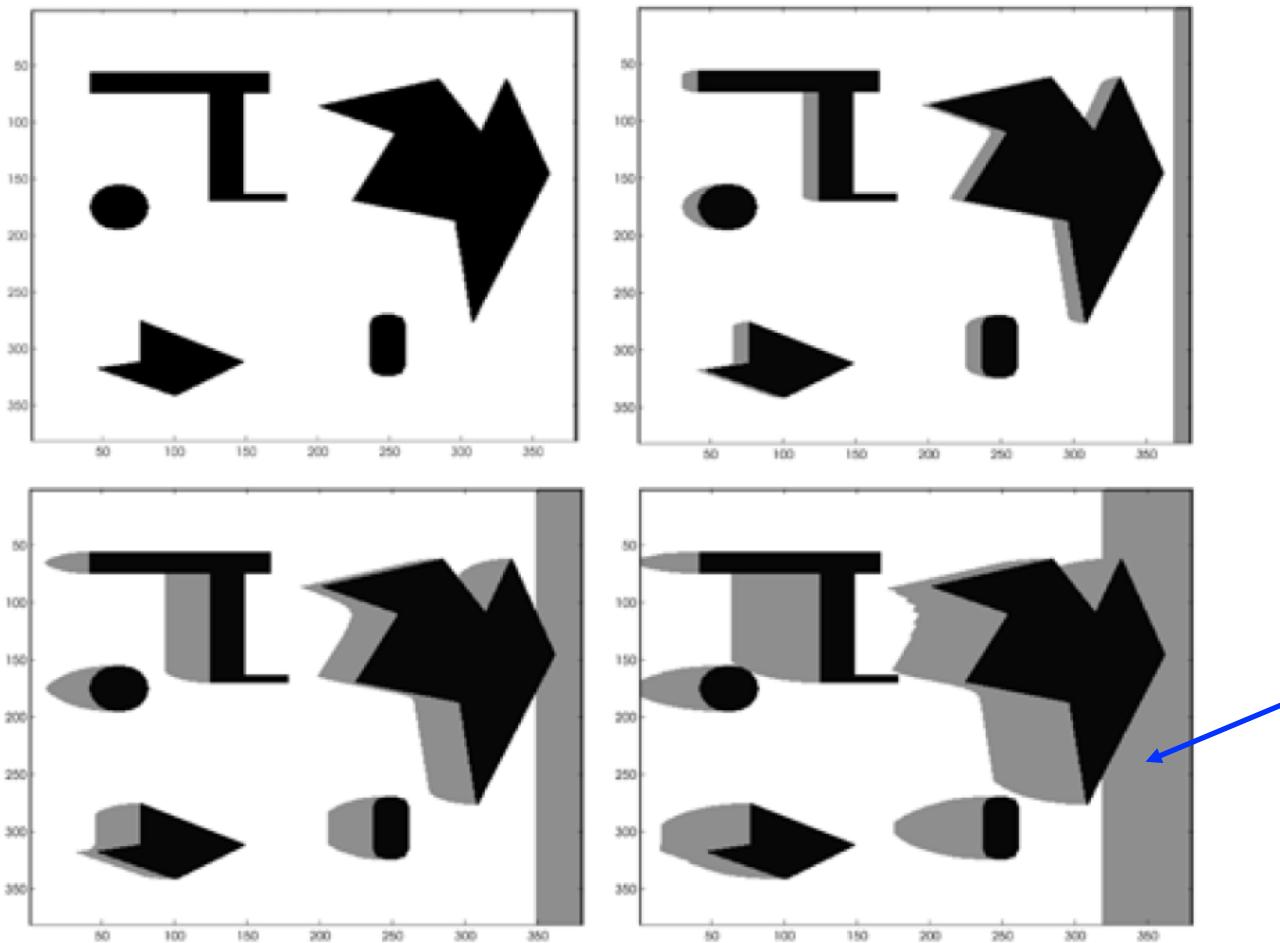
- Probabilistically complete
- Single query & directed
 - Edges in RRTs are **directed**: travel is **from** the initial location **towards** the goal ($x \rightarrow y \neq y \rightarrow x$)
 - RRTs **retain control information u** in their edges
 - **Key payoff:** Unlike PRMs, RRTs can easily handle **dynamic constraints**:

$$\dot{x} = f(x, u)$$

In fact, RRTs were explicitly developed for ***kinodynamic planning*** (planning w/
kinematic and dynamic constraints)

⇒ Often applied to planning in ***phase space*** (position *and* velocity)

Planning with phase space as configuration space

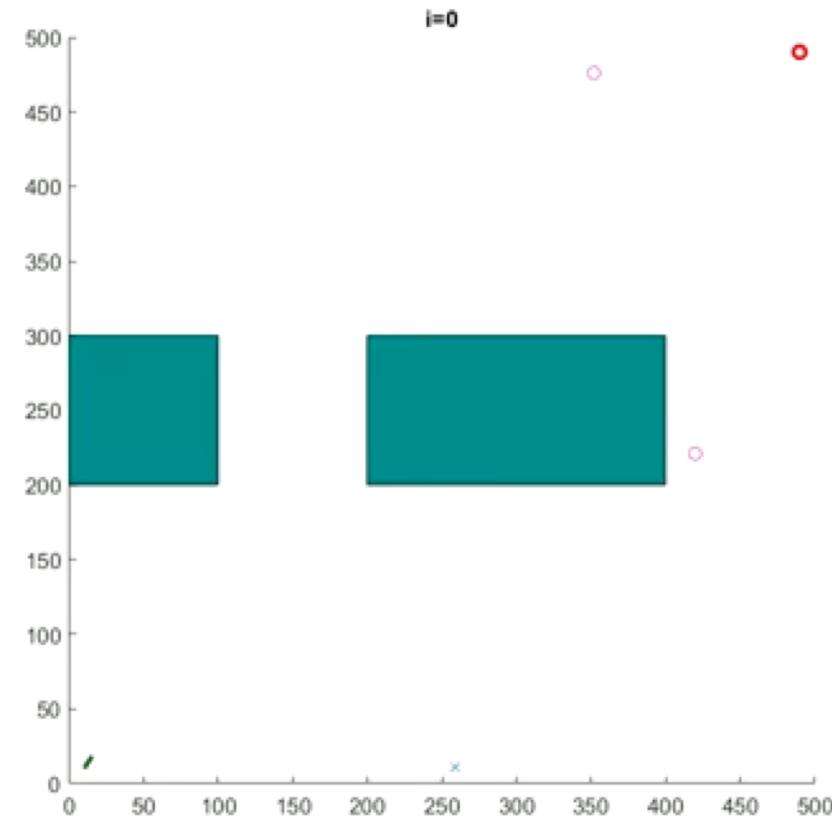
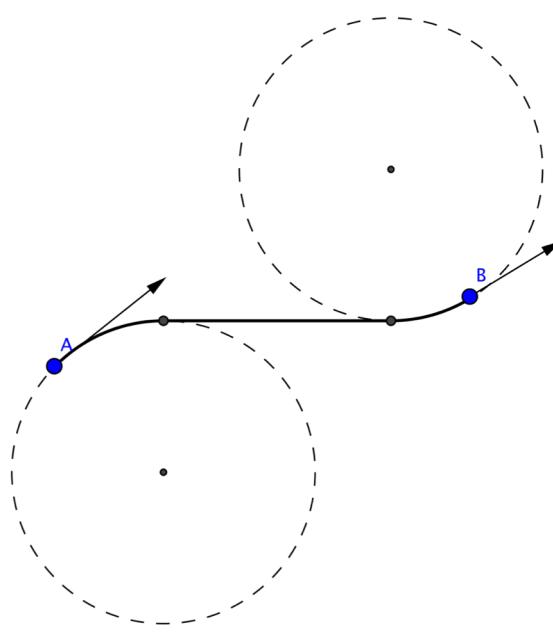


Region of inevitable collision grows with respect to the robot's initial speed.

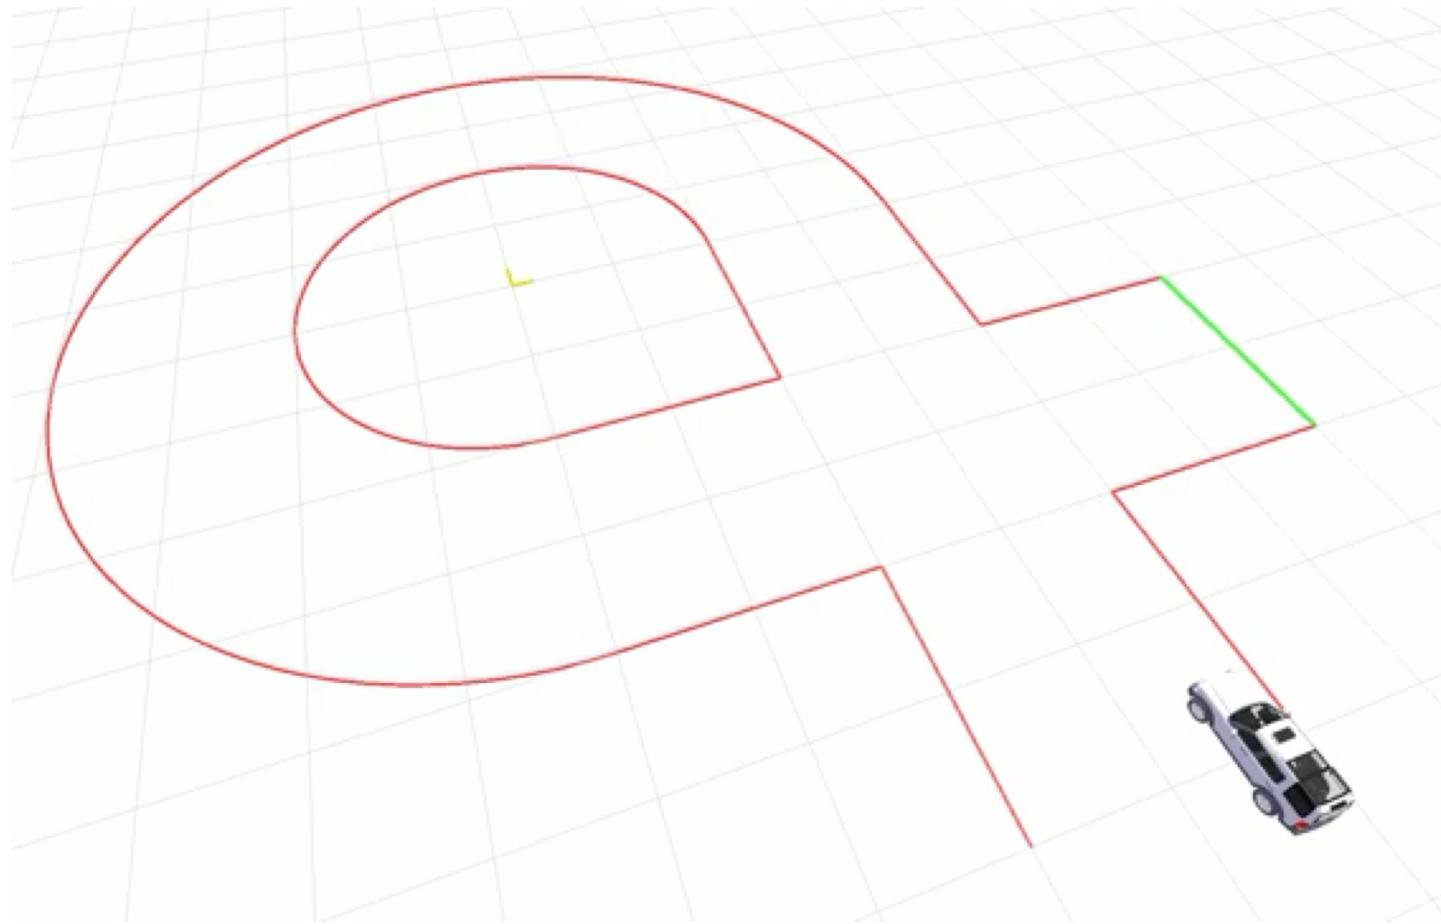
Fig. 2. Slices of \mathcal{X} for a point mass robot in two dimensions with increasingly higher initial speeds. White areas represent \mathcal{X}_{free} , black areas are \mathcal{X}_{obst} , and gray areas approximate \mathcal{X}_{ric} .

Example: RRT Dubin's path planner

$$\begin{aligned}\dot{x} &= V \cos(\theta) \\ \dot{y} &= V \sin(\theta) \\ \dot{\theta} &= u\end{aligned}$$



RRT trajectory planning with racecar dynamics



Other variants of RRT

- Various sampling ideas
 - Finding goal quickly
 - Sampling from narrow passages
- Construction of the tree in two directions
- RRT* (rewiring and cost incorporation) – **asymptotically optimal**

Variations on a Theme: Biased RRT sampling

Input: q_{start} , q_{goal} , number n of nodes, stepsize α , β

Output: tree $T = (V, E)$

```
1: initialize  $V = \{q_{\text{start}}\}$ ,  $E = \emptyset$ 
2: for  $i = 0 : n$  do
3:   if  $\text{rand}(0, 1) < \beta$  then  $q_{\text{target}} \leftarrow q_{\text{goal}}$ 
4:   else  $q_{\text{target}} \leftarrow \text{random sample from } Q$ 
5:    $q_{\text{near}} \leftarrow \text{nearest neighbor of } q_{\text{target}} \text{ in } V$ 
6:    $q_{\text{new}} \leftarrow q_{\text{near}} + \frac{\alpha}{|q_{\text{target}} - q_{\text{near}}|} (q_{\text{target}} - q_{\text{near}})$ 
7:   if  $q_{\text{new}} \in Q_{\text{free}}$  then  $V \leftarrow V \cup \{q_{\text{new}}\}$ ,  $E \leftarrow E \cup \{(q_{\text{near}}, q_{\text{new}})\}$ 
8: end for
```

Biassing tree expansion towards the goal: Sample the goal state q_{goal} itself with probability $\beta > 0$

Variations on a Theme: Bi-directional RRT

```
RRT_BIDIRECTIONAL( $x_{init}, x_{goal}$ );
1    $\mathcal{T}_a.init(x_{init}); \mathcal{T}_b.init(x_{goal});$ 
2   for  $k = 1$  to  $K$  do
3        $x_{rand} \leftarrow \text{RANDOM\_STATE}();$ 
4       if not ( $\text{EXTEND}(\mathcal{T}_a, x_{rand}) = \text{Trapped}$ ) then
5           if ( $\text{EXTEND}(\mathcal{T}_b, x_{new}) = \text{Reached}$ ) then
6               Return PATH( $\mathcal{T}_a, \mathcal{T}_b$ );
7               SWAP( $\mathcal{T}_a, \mathcal{T}_b$ );
8       Return Failure
```

Fig. 7. A bidirectional rapidly exploring random trees-based planner.

Bi-directional RRT grows two trees towards each other: one from the initial state, and one from the goal

Example: Bi-directional RRT path-planning

